

Advanced Algorithms

Lecture 9

External-Memory Algorithms and Data Structures

Bin Yang

byang@cs.aau.dk

Center for Data-intensive Systems

Intended Learning Outcomes (ILO)



- After taking this course, you should acquire the following knowledge
 - Algorithm **design** techniques such as **divide-and-conquer**, **greedy algorithms**, **dynamic programming**, back-tracking, branch-and-bound algorithms, and **plane-sweep algorithms**;
 - Algorithm **analysis** techniques such as **recursion**, **amortized analysis**;
 - A collection of **core** algorithms and data structures to solve a number problems from various computer science areas: *algorithms for external memory*, multiple-threaded algorithms, **advanced graph algorithms**, heuristic search and **geometric calculations**;
 - There will also enter into one or more **optional subjects** in advanced algorithms, including, but not limited to: *approximate algorithms*, randomized algorithms, search for text, linear programming and number theoretic algorithms such as cryptosystems.

ILO of Lecture 9



- External memory algorithms and data structures
 - to understand the external memory model and the principles of analysis of algorithms and data structures in this model;
 - to understand the algorithms of B-tree and its variants and to be able to analyze the complexity;
 - to understand the main principles of external tree structures;
 - to understand how the different versions of merge-sort algorithms work in external memory;
 - to understand why the amount of available main-memory is an important parameter for the efficiency of external-memory algorithms.

Agenda




- External memory
- B-trees, B⁺-trees, and R-trees
- External memory merge sort

Hard disks, magnetic disks




- In real systems, we need to cope with data that does not fit in main memory.
- How does a hard disk work:
 - It has one or more *platters* that rotate around a *spindle*.
 - Each platter is read/write with a head at the end of an arm.
 - Arms rotate around a common pivot axis.
 - Track is the surface under the head.



Hard disks, magnetic disks



- Reading data from the hard disk:
 - *Seek* with the head.
 - *Wait* while the necessary sector rotates under the head
 - *Transfer* the data.



Some numbers




- Rotation speed: 5,400 to 15,000 RPM.
 - Laptop hard disk, normally 5,400 or 7,200.
 - My laptop's hard disk: WD5000LPVX: 5,400
- Let's consider a hard disk with 7,200 RPM
 - One rotation takes $60/7200 \approx 8.3$ milliseconds (10^{-3})
 - Plus some time for moving the arms.
 - Main memory: 50 nanoseconds (10^{-9})
 - The hard disk takes 5 orders of magnitude longer access times.
- Conclusions for using hard disks
 - Disk access is much slower than main-memory access
 - Thus, it makes sense to read and write in large blocks – *disk pages* (4 – 32Kb)
 - *Sequential* access is much faster than *random* access

SSDs



- The same, although to less extent is true for *flash-based solid state drives* (SSDs):
 - It is more efficient to read/write (especially write) in larger blocks
 - Sequential/random I/O difference is less pronounced than in disks.
- Depth of the memory hierarchy (access latency):
 - DRAM(~50ns) – x2000 → SSD(~0.1ms) – x50 → HDD(~5ms)

Historical Cost of Computer Memory and Storage



Picture from hblok.net/blog/storage

External memory model



- Two principal components of the running time analysis:
 - Not only, the CPU time, i.e., the computing time.
 - But also, and more importantly, the number of disk page accesses, i.e., the number of I/O.
- **B** – page size is an important parameter:
 - *Example:* $n = 256\text{MB}$, $B = 4\text{KB}$, 0.1 ms disk access
 - ◆ n disk accesses: $256 \cdot 10^6 \cdot 0.1 \cdot 10^{-3} = 25\,600\text{s} = 7.1\text{ hours}$
 - ◆ n/B disk accesses: $= 64 \cdot 10^3 \cdot 0.1 \cdot 10^{-3} = 6.4\text{s}$
 - ◆ Read/Write in blocks but not individual data instances
 - Not “just” a constant:
 - ◆ $\Theta(\log_2 n) \neq \Theta(\log_B n)$
 - ◆ $\Theta(n) \neq \Theta(n/B)$

External memory algorithms



- The typical working pattern for algorithms:

```
01 ...
02  $x \leftarrow$  a pointer to some object
03 DiskRead(x)
04 operations that access and/or modify x
05 DiskWrite(x) //omitted if nothing changed
06 other operations, only access no modify
07 ...
```

- If the object referred to by x resides on disk, **DiskRead**(x) needs to read object x into main memory before we can access or modify x .
- If the object referred to by x is already in main memory, **DiskRead**(x) does nothing and does not incur another time disk access.

Agenda




- External memory
- B-trees, B⁺-trees, and R-trees
- External memory merge sort

B-trees



- B-tree is a balanced tree.
 - All leaf nodes have the same depth.
- Internal nodes
 - x multiple keys and $x+1$ pointers to child nodes.
 - $\text{pointer}_1 \text{ key}_1 \text{ pointer}_2 \text{ key}_2 \text{ pointer}_3 \text{ key}_3 \dots \text{pointer}_x \text{ key}_x \text{ pointer}_{x+1}$
 - $\text{key}_1 \leq \text{key}_2 \leq \text{key}_3 \leq \dots \leq \text{key}_x$
 - For the first and last pointers: $\text{pointer}_1.\text{key} \leq \text{key}_1$
and $\text{key}_x < \text{pointer}_{x+1}.\text{key}$
 - For the remaining pointers: $\text{key}_{i-1} < \text{pointer}_i.\text{key} \leq \text{key}_i$
- Leaf nodes
 - Only keys.




B-trees



- Each node resides on a disk page.
- Given a page size **B**, how many keys can a page hold at most?
 - Assume that a key is a 4 bytes integer and let a pointer be 8 bytes.
 - For a $B=4k$ page setup, we have $4m+8(m+1) \leq 4096$. Then, $m=340$.
- Nodes have **lower** and **upper** bounds on the number of keys they can contain
 - Upper bound: each node has at most m keys and $m+1$ children (max_fan-out)
 - Lower bound: each node has at least $t = \lfloor m/2 \rfloor$ keys and $t+1$ children (min_fan-out)
 - Example:
 - ◆ $m=5$: at most 5 keys 6 children, and at least $t=2$ keys 3 children.
 - ◆ $m=6$: at most 6 keys 7 children, and at least $t=3$ keys 4 children.
- Root is an exception: root node can have as little as only one key and two children.

B-trees



- Each node resides on a page.
- $m=4$, $t = \lfloor m/2 \rfloor = 2$
 - Each node has at least 2 keys and at most 4 keys.
 - Each node has at least 3 children and at most 5 children.
 - Root is an exception, only has 1 key and 2 children.

Searching on B-trees



- The root node is normally “always” in main memory.
 - No need to perform a DiskRead on the root.
- Search is very similar to a search in a binary search tree
 - Instead of making a binary branching decision at each node, we make a $(j+1)$ -way branching decision, where j is the number of keys in a node.

Pseudo code



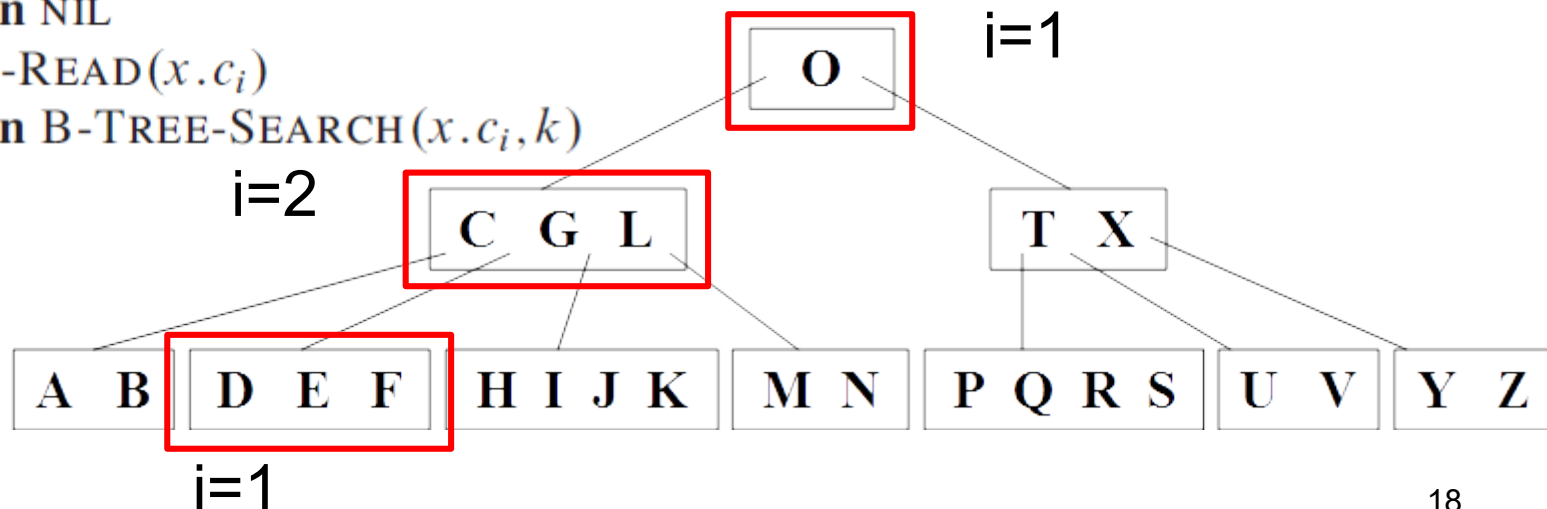
- x is a node and $x.n$ is the number of keys in the node.
- k is the key that we are searching for.
- $x.key_i$ is the i -th key of node x ; and $x.c_i$ is the i -th pointer of node x .

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return ( $x, i$ )
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

Searching for “D”, i.e., $k = D$
B-Tree-Search(root, D)


Disk access: $O(h) = O(\log_t n)$
CPU: $O(th) = O(t \log_t n)$



B⁺-trees




- B⁺-trees is a variant of B-trees:
 - All data keys are in leaf nodes.
 - Leaf nodes are connected into a linked list.
 - Extensively used in database systems for indexing data records that are stored in a database.



Searching on B⁺-trees



- Searching
 - Always goes to a leaf
 - ◆ O in B-trees vs. Θ in B⁺-trees
 - Point search: $\Theta(\log_t n)$
 - Range search: $\Theta(\log_t n + k/t)$



B⁺-trees: Insertion




- Skeleton of the algorithm:
 - *Down-phase*: recursively traverse down and find the leaf (as in search)
 - ◆ If the leaf has room, just insert it in the leaf. Otherwise:
 - *Up-phase*: *split* nodes and propagate the splits up the tree
 - ◆ Split the leaf into two leaves and divide the keys between the two new nodes. **Copy** the middle key to its parent.
 - ◆ Apply the same strategy to insert the middle key to its parent.
 - If there is room, insert directly. Otherwise, split and **move** the middle key to its parent.
 - ◆ When the root is full, then we create a new root with a single key.


B⁺-trees: node splitting



- Assume that at most 6 keys in a node.
- Leaf node (**copy** the middle key to the parent)



- Internal node (**move** the middle key to the parent)

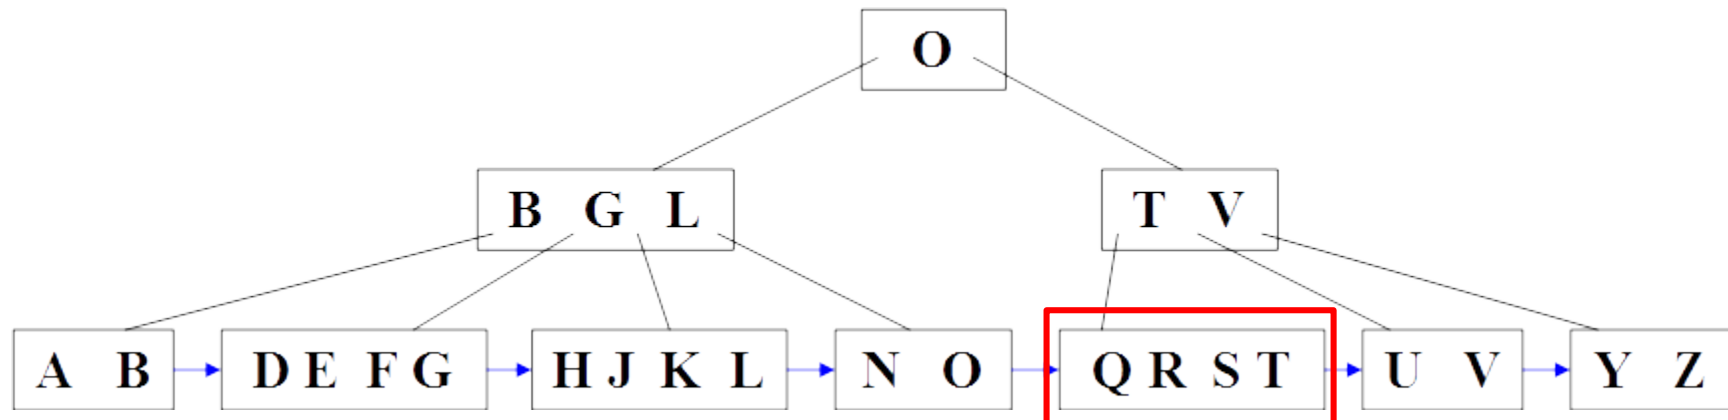


- The tree grows when the root is split into two nodes and their parent becomes the new root.


Example



- *Insert P* (assume that at most 4 keys)




Down-phase: Leaf node
(**copy** the middle key to the parent)




Mini-quiz (also on Moodle)




- At most 4 keys
- After inserting I, C
- How does the root node look like?






- Insert I
 - HJKL becomes HIJKL. Then split HIJ, KL. J is copied to its parent.

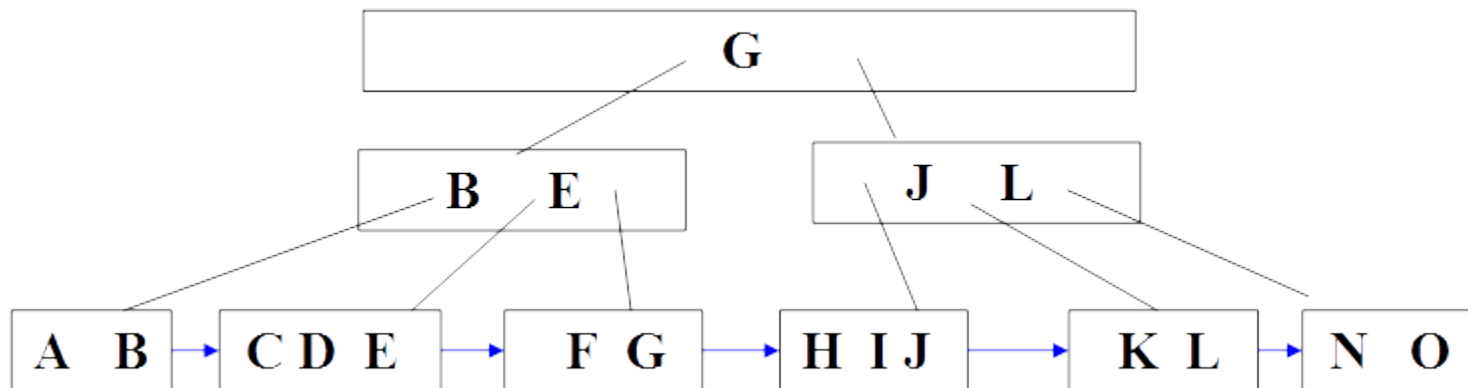




- Insert C
 - DEFG becomes CDEFG. Split CDE, FG. E is copied to its parent.



- BEGJL splits to BE and JL and a new root with G is created.



B⁺-trees: Deletion



- Opposite of insertion:
 - Down-Phase: traverse down to find the key in a leaf
 - Up-Phase: remove the key and traverse up handling underfull nodes
- Tree shrinking: if the root has only one child, remove the root.

B⁺-trees: Deletion (under-full)

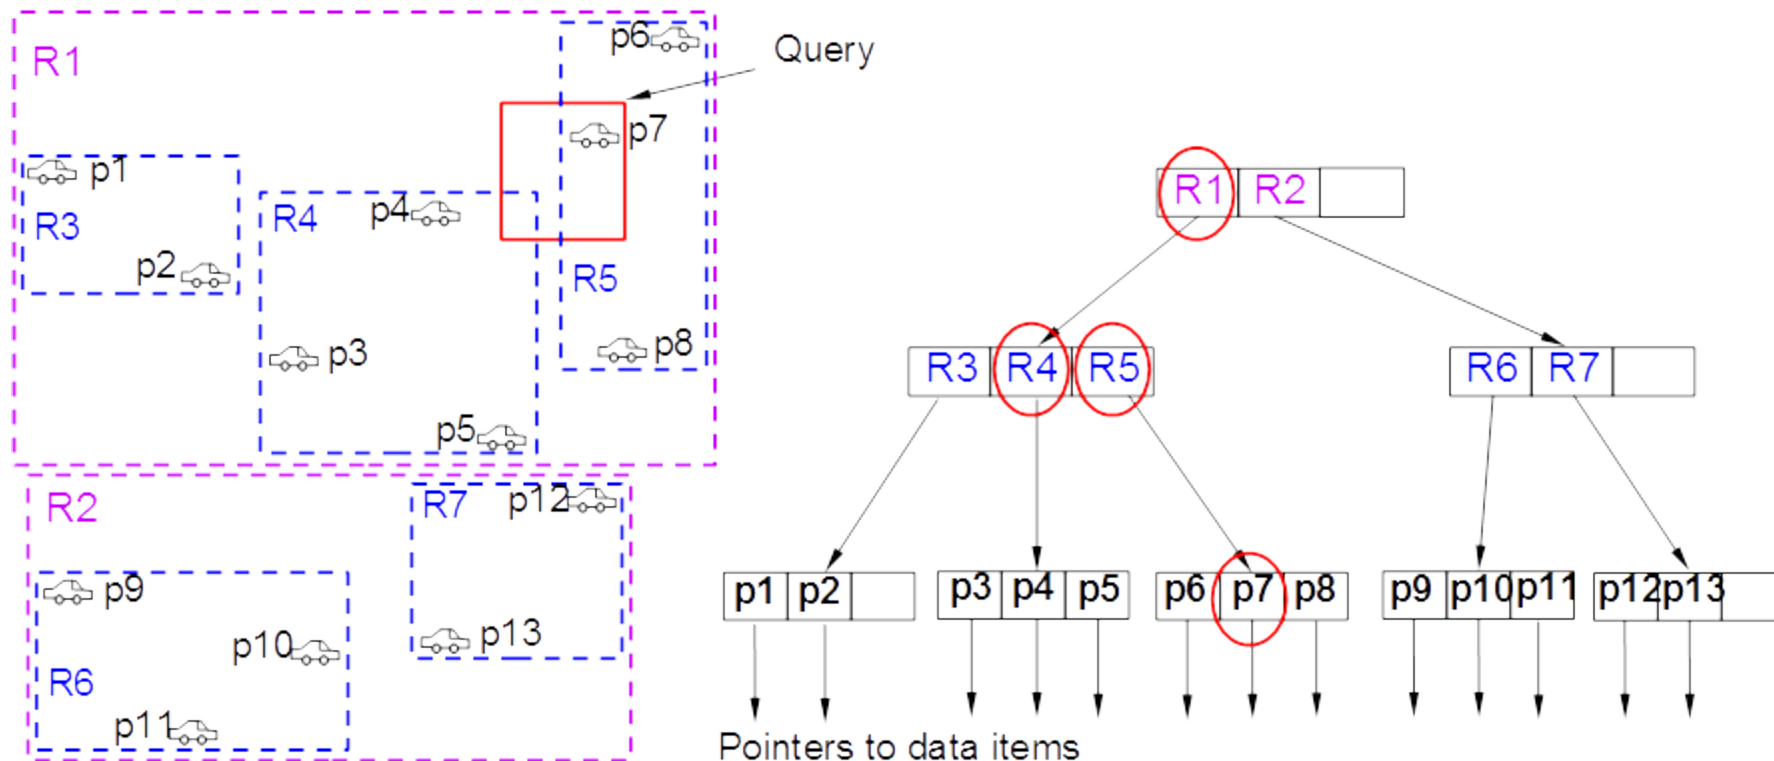


- After deleting a key from node N, N becomes underful, i.e., less than the minimum number of keys/children.
- Case 1: If one of the adjacent sibling nodes of N has more than the minimum number of keys, move one from its sibling node to N.
- Case 2: If neither adjacent sibling node can provide the extra key. Merge N and one of its adjacent sibling into one node.
 - One node has less the minimum number of keys, and one has exactly the minimum number of keys. Combining them into one node won't exceed the maximum number of keys.
 - This is why we set the minimum number of keys to be half of the maximum number of keys.

R-trees



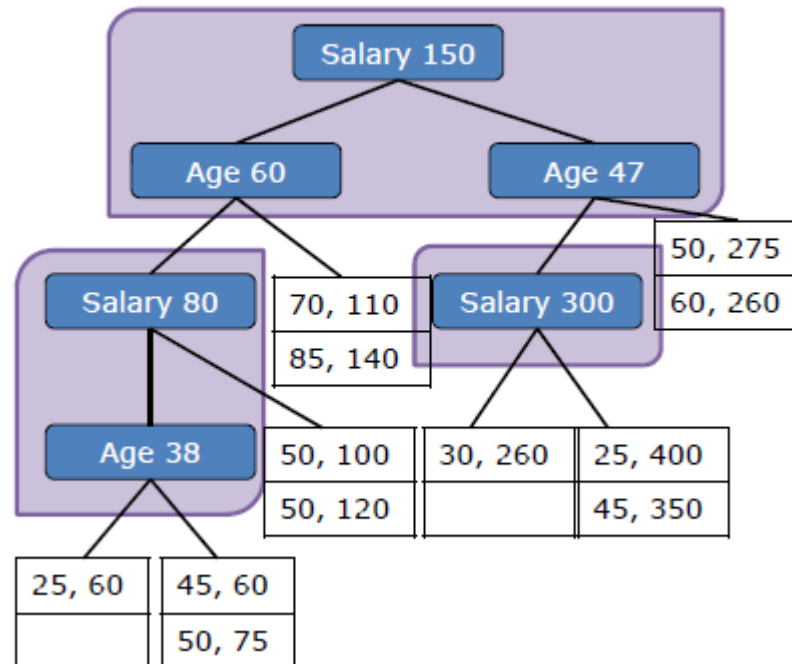
- A multi-dimensional extension of B-trees.
- Key is not a value, but a multi-dimensional range.
 - 2D: A rectangle
- Mini-quiz: how many rectangles can a node hold?



What about kd-trees?





- Since inner node of a kd-tree is tiny, we may group several inner nodes and store them in one disk page.
 - To minimize the number of disk pages that we must read from disk while traveling down one path, it is good to group a node with its children inner nodes for some number of levels.






- Combining the characters of kd-trees and B-trees
 - Multidimensional search efficiency of kd trees.
 - I/O efficiency of B-trees, multi-way branch, balanced tree.
 - ◆ Fan out determined by page size and size of entries.
- Region nodes (internal nodes)
 - (region, pointer-to-a-child-node) pairs.
 - Region is defined as min/max per dimension.
 - Regions are disjoint, and union of all regions in a node is a region.
 - Region of the root is the whole space.
- Point nodes (leaf nodes)
 - (point, pointer-to-data-record) pairs.
- Nodes=pages, each node is stored on a disk page.





Point
query



Design principles of external memory trees



- What do we want?
 - A short tree, because it decides how many pages will be loaded, i.e., the number of I/Os.
- How to achieve this?
 - Fully utilize a page to have many entries in a node.
- Short, but wide, trees.

Agenda



- External memory
- B-trees, B⁺-trees, and R-trees
- External memory merge sort

External-Memory Sorting



- External-memory algorithms
 - When data do not fit in main-memory
- External-memory sorting
 - Rough idea: sort pieces that fit in main-memory and “merge” them
- Main-memory merge sort:
 - The main part of the algorithm is Merge
 - Let's merge:
 - ◆ 3, 6, 7, 11, 13
 - ◆ 1, 5, 8, 9, 10

Main-memory merge sorting



Merge-Sort(A)

01 **if** length(A) > 1 **then**

02 Copy the first half of A into array A1

03 Copy the second half of A into array A2

04 **Merge-Sort**(A1)

05 **Merge-Sort**(A2)

06 **Merge**(A, A1, A2)

} *Divide*
} *Conquer*
} *Combine*

Running time?

Recurrence: $T(n) = 2T(n/2) + \theta(n)$
 $\theta(n \lg n)$

External-memory merge sort



- Settings
 - The total number of input elements that we want to sort is N .
 - Available memory can hold M elements for in-memory merge sort.
 - A disk page can hold B elements.
 - $N > M > B$.
 - $n = N/B$: the total number of disk pages of the input.
 - $m = M/B$: the total number of disk pages that fit in the available memory.
- Example: $N = 10^6 > M = 10^3 > B = 10^2$
 - $n = 10^4$ disk pages of elements in total to be sorted.
 - $m = 10$ disk pages can fit in the available memory.

External-memory merge sort



- Input file X , empty file Y with the same size of X .

- *Phase 1*: Repeat until the end of file X :

- Read the next M elements from X .
- Sort them in main-memory.
 - ◆ We call them a run, i.e., a sorted sub-array.
- Write them at the end of file Y .
- *In total, N/M runs, each with M sorted elements.*

What is the complexity w.r.t. disk page visits (i.e., # of I/O) in the first phase?

- Assume that: $N=8000$, $M=1000$, $B=25$


$$2n=2 \cdot 320$$

- $n=320$ pages, $m=40$ pages

- After phase 1, we have 8 runs in file Y .

- For each run, we have 1000 elements which are sorted.

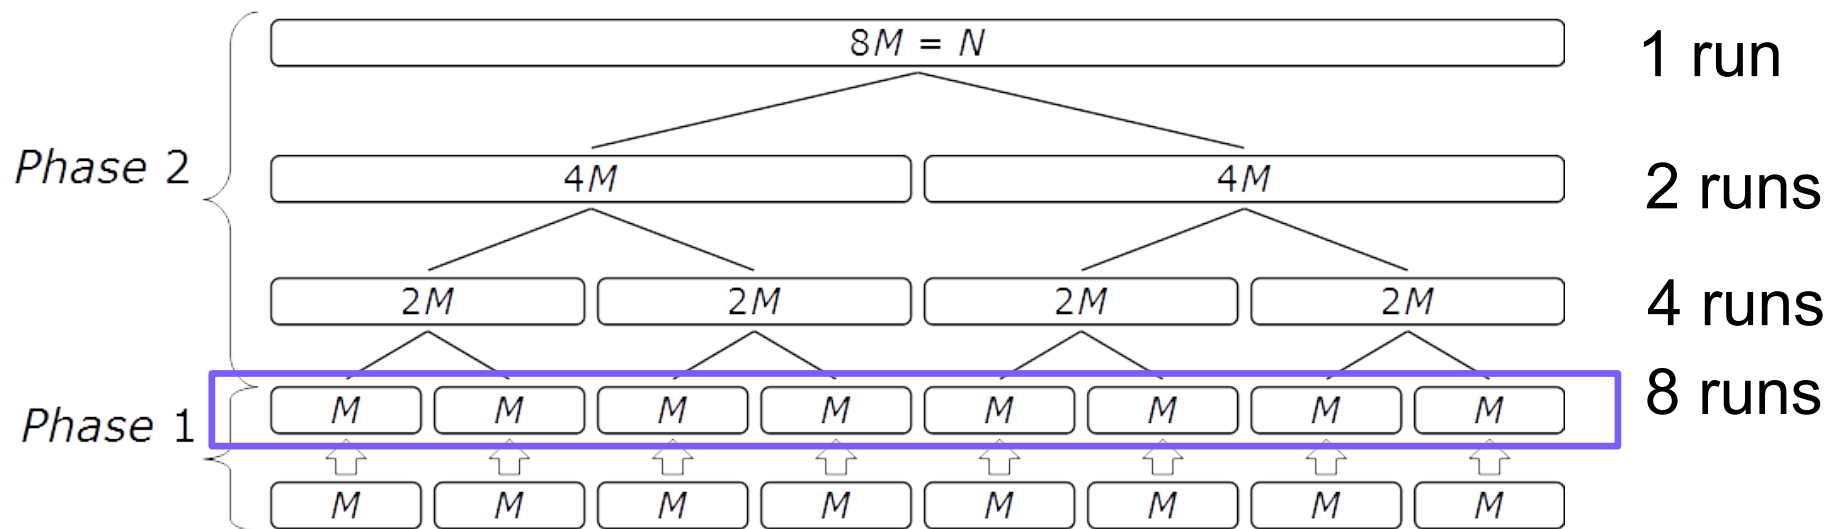
Y : 8 runs. Each run is sorted already.



External-memory merge sort



- *Phase 2*: Repeat while there is more than one run in Y :
 - Empty X
 - *MergeAllRuns* (Y, X)
 - X is now called Y , Y is now called X



External two-way merge

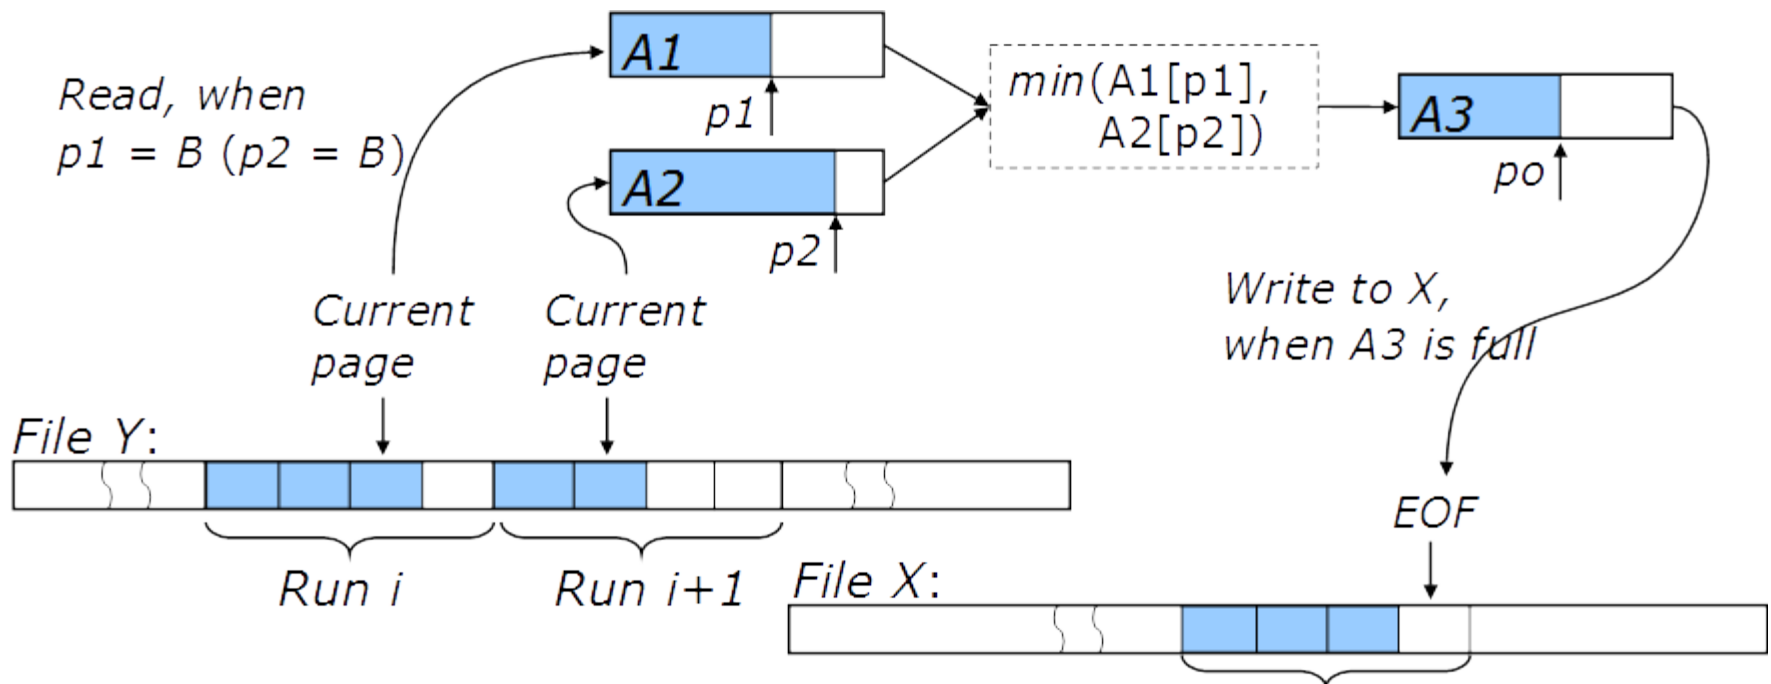


- In Phase 2, we cannot use Merge() function from the main-memory merge sorting
 - Because each run has M elements already, and we cannot load both runs into the main memory to perform the Merge() function.
- Instead, we use **two-way** merge to merge **two** runs into **one longer** run.
 - We use three main-memory arrays of size B , i.e., 3 main-memory pages
 - ◆ A_1, A_2, A_3 .
 - Copy one page from the first run to A_1
 - Copy one page from the second run to A_2 .
 - Merge A_1 and A_2 to A_3 using main-memory Merge.
 - Whenever A_3 is full, store A_3 to file X .

External two-way merge



- Every time, merge two runs into one longer run.
 - ◆ Copy one page from the first run to A1
 - ◆ Copy one page from the second run to A2.
 - ◆ Merge A1 and A2 to A3 using main-memory merge.
 - ◆ Store A3 to file X when A3 is full.



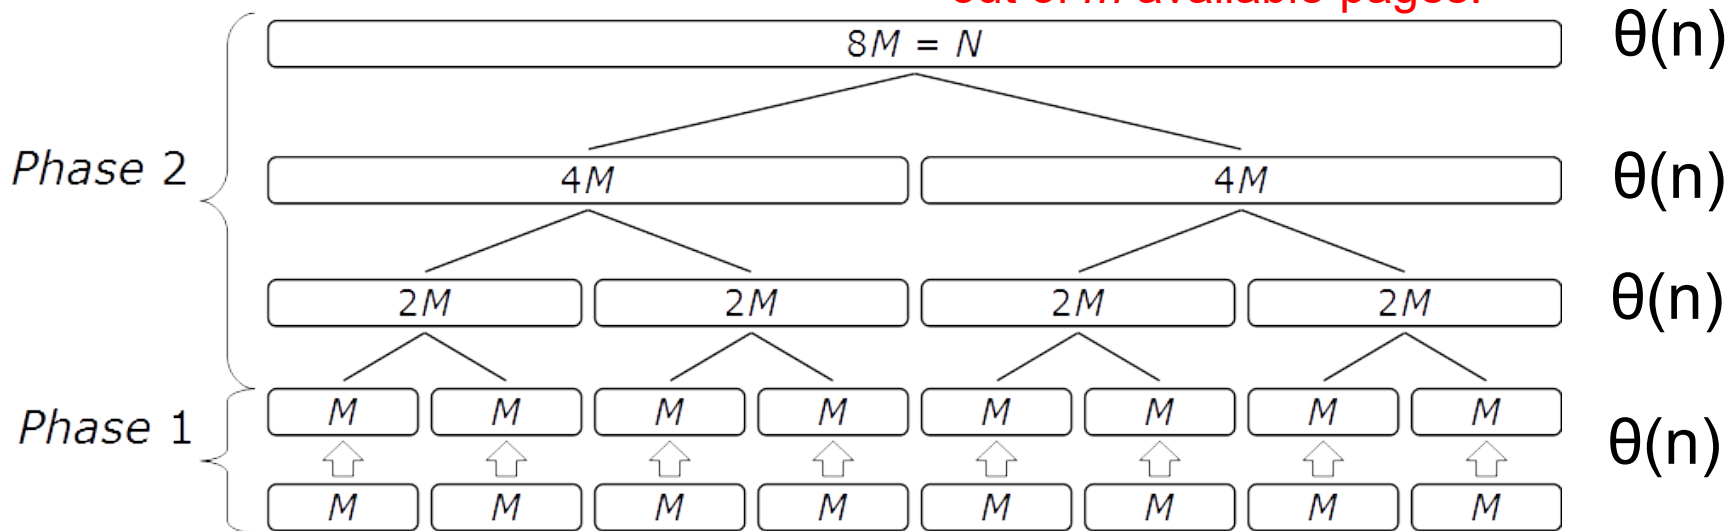
Analysis



- Phase 1: $n=N/B$ reads and $n=N/B$ writes. Thus $\theta(n)$.
- Phase 2:
 - Each iteration: n reads and n writes. Thus $\theta(n)$.
 - How many iterations:
 - ◆ We start with $\lceil N/M \rceil = \lceil n/m \rceil$ runs and stop with only 1 run.
 - ◆ Each iteration we reduce the number of runs by half.
 - ◆ $\theta(\lg(n/m))$ iterations in total.
 - Thus, $\theta(n \lg(n/m))$ read/write.

We can do better!

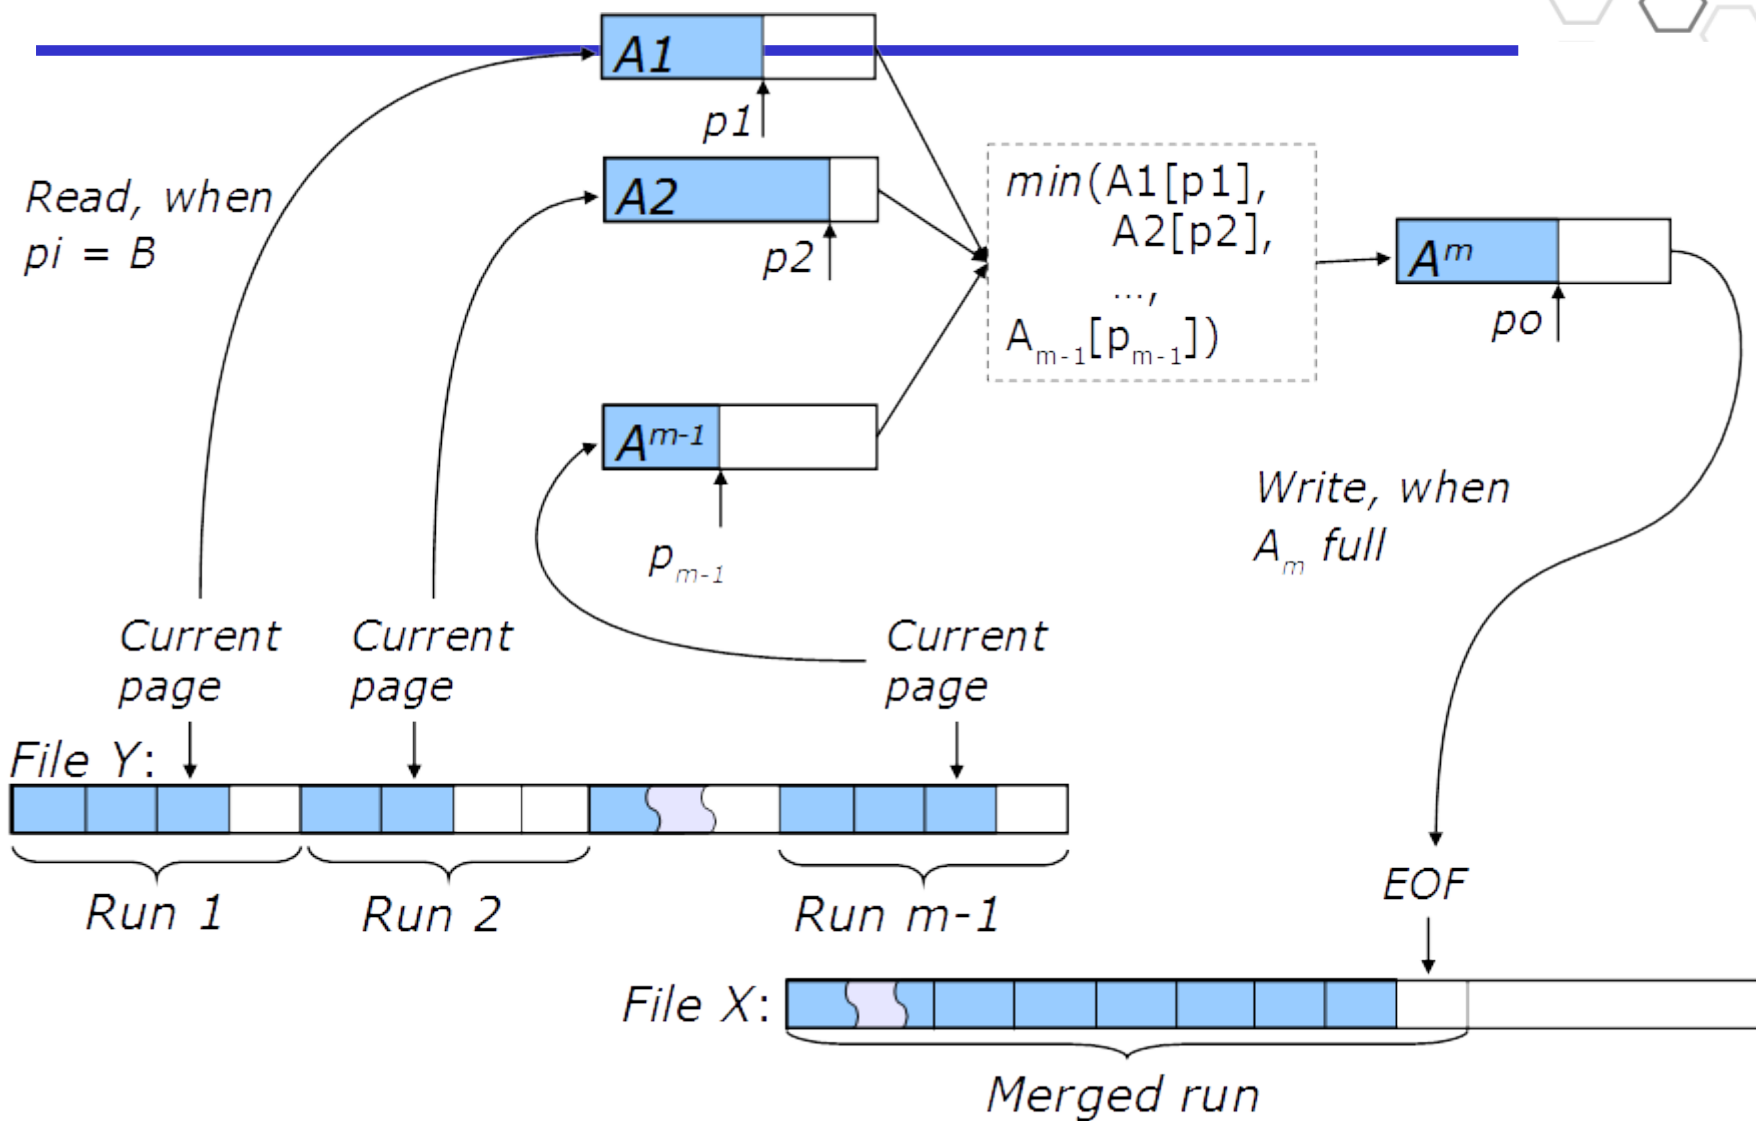
Observation: Although phase 1 uses all available memory, phase 2 uses just 3 pages out of m available pages!



External multi-way merge sort



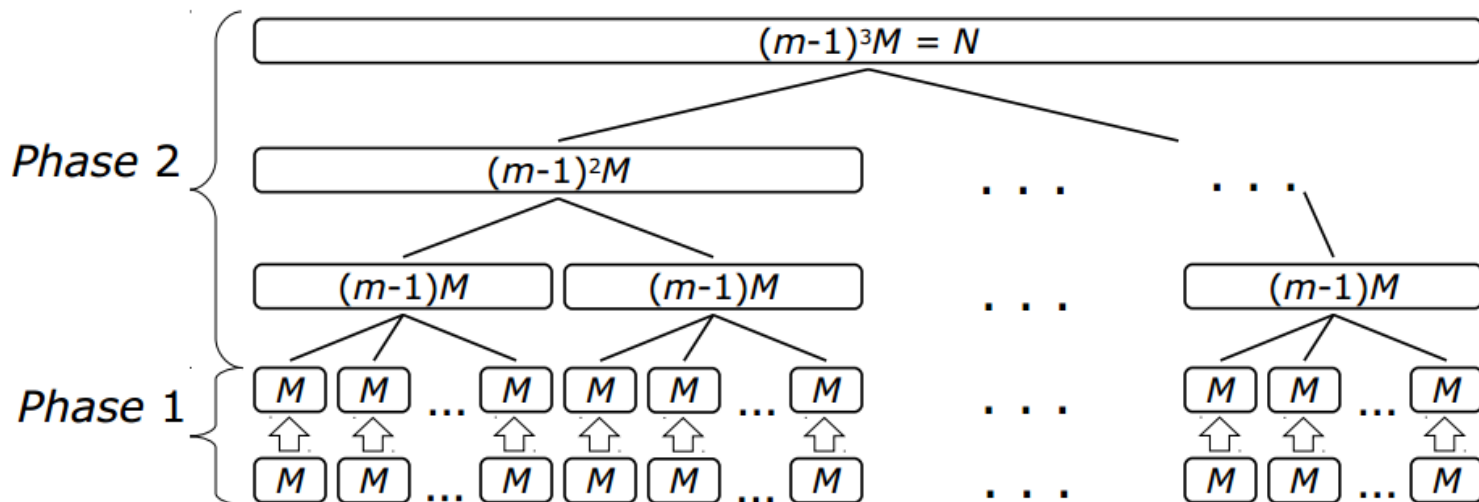
- In phase 2, we use **multi-way** merge to merge **$m-1$** runs into **one longer** run.
 - Instead of using **only three** main-memory pages of size B , we use **all m** main-memory pages of size B .
 - Copy one page from the first run to A_1
 - Copy one page from the second run to A_2 .
 - Copy one page from the third run to A_3 .
 - ...
 - Copy one page from the $(m-1)$ -th run to A_{m-1} .
 - Merge $A_1 A_2, \dots, A_{m-1}$ to A_m using main-memory merge.
 - Whenever A_m is full, store A_m to file X .
- In our running example, we can afford merging 39 runs into one longer run in one iteration.
 - We only have 8 runs in total. So only one iteration is needed.
 - In contrast, we need 3 iterations when using two-way merge sort.



External multi-way merge sort



- In most cases, if the input file is **not extremely large**, we only need to do phase 2 only once.
 - Refer to the reading materials on Moodle to get an idea about what is “extremely large”.
 - $\theta(n)$
- If the input file is extremely large, we have to do multi-way merge $\log_{m-1}(n/m) = \theta(\log_m n)$ times.
 - $\theta(n \log_m n)$





- External memory algorithms and data structures
 - to understand the external memory model and the principles of analysis of algorithms and data structures in this model;
 - to understand the algorithms of B-tree and its variants and to be able to analyze them;
 - to understand the main principles of external tree structures;
 - to understand how the different versions of merge-sort derived algorithms work in external memory;
 - to understand why the amount of available main-memory is an important parameter for the efficiency of external-memory algorithms.