

Advanced Algorithms

Lecture 1 *Introduction* *&* *Dynamic Programming*

Center for Data-intensive Systems

Bin Yang

byang@cs.aau.dk

Agenda



- Introduction
- Dynamic Programming

People



- Lecturers: Bin Yang, Tung Kieu, Chenjuan Guo
 - Email: byang@cs.aau.dk
 - Office: 3.2.46
 - Homepage: <http://people.cs.aau.dk/~byang/>
 - DPW group: **Database**, Programming and Web Technologies
 - ◆ Big data, data science, data analytics (machine learning, artificial intelligence)
 - Daisy: Center for Data-Intensive Systems.
<http://www.daisy.aau.dk/>

Location and Time




- Location
 - 0.1.95.
- Time
 - Tuesdays and Fridays.
 - Lectures are from 8.15 to 10.00.
 - Exercises are from 10.15 to 12.00.
- Check the full schedule on Moodle in case there are last-minute changes.

Moodle Page and Textbook



- Course page at Moodle
 - <https://www.moodle.aau.dk/course/view.php?id=33044>
 - Please check it frequently for notifications and updates!
- Textbook
 - “Introduction to Algorithms”, 3.ed, by Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein, The MIT Press.
 - ISBN:9780262533058
 - Abbreviation: CLRS
 - Available at the Factum book store <http://ftu.dk/>



Course Structure



- A total of 15 sessions
 - 12 regular sessions + 3 self-study exercise sessions
- A regular session = a lecture class + an exercise class.
- A lecture class
 - 2 * 45 minutes
 - Each lecture we focus on a specific topic.
 - During a lecture, there may be mini-quizzes.
 - ◆ Have a pen and some papers.
- An exercise class
 - 2 * 45 minutes
 - Solve exercises assigned in the session.
 - ◆ You are encouraged to work in groups.
 - ◆ TAs come by from group to group during each exercise class.
 - Get a feeling of exam questions from enough exercises!

Course Structure (2)



- A self-study exercise session = 4 hours of exercises
 - You need to do it in groups.
 - Each group **can** submit to TA or me one written solution no later than a week of the session.
 - TA/I will give written feedback for each of the submitted solutions.
 - I recommend that each of you solves the problems **individually** first, and then you discuss, summarize, and hand in solutions **per group**.
 - In case you cannot agree with each other, you can hand in multiple solutions for one problem.

Working hours



- This is a 5 ECTS course, where each ECTS=27.5 hours
 - 137.5 hours
- 12 regular sessions
 - 2h lecture + 2h exercises.
 - 3h reading
 - In total, $12 \times (2+2+3) = 84\text{h}$
- 3 self-study exercises sessions
 - 4h on solving the exercises.
 - 4h on checking the solutions/feedback in order to make sure you can solve each of the exercises.
 - In total, $3 \times (4+4) = 24\text{h}$
- 29.5h for preparing the exam.
- In total, $84+24+29.5=137.5\text{h}$

Exam



- Individual and written, but open-book
 - Electronic devices with communication capabilities, such as laptops and mobile phones, are NOT allowed.
 - You can bring old-fashion calculators.
 - You can freely use your copies of slides from the lectures, textbooks, and other course material.
- Exam in two parts
 - A set of quizzes, to test knowledge.
 - ◆ Choose the options that you think are correct.
 - A few (1 ~ 3) open problems, each with some sub-problems, to test competences and skills.
 - ◆ Given a real world problem, write pseudo code, give complexity analysis, etc.
 - You will get a better feeling when you participate self-study exercise sessions.
- The exam will have 100 points.

Prerequisites



- AD on DAT3/SW3 or AD2 on IT7
- Let's quickly recap the intended learning outcomes of AD/AD2
- Basic mathematical concepts such as recursion, induction, concrete and abstract complexity;
 - Solving recurrences, asymptotic notation.
- Basic data structures;
 - Queues, stacks, heaps, linked lists, priority queues.
- Algorithmic principles such as searching, search trees, sorting, dynamic programming, divide-and-conquer;
 - Binary search tree, merge sort, quick sort.
- Graphs and graph algorithms such as shortest path, connected components, spanning trees.
 - BFS, DFS, MST, topological sorting, shortest path.

Mini quiz



1.2. (3 points) $700 \cdot n^2 + 999 \cdot n^2 \lg n + 0.1 \cdot n^2 \lg^2 n$ is:

- ☐ a) $\Theta(n^2 \lg n)$ ☐ b) $\Omega(n^2 \lg n)$ ☐ c) $\Theta(n^2)$ ☐ d) $\Theta(n^2 \cdot \lg^2 n)$

Intended Learning Outcomes (ILO)



- After taking AALG, you should acquire the following knowledge
 - Algorithm **design** techniques such as divide-and-conquer, greedy algorithms, dynamic programming, back-tracking, branch-and-bound algorithms, and plane-sweep algorithms;
 - Algorithm **analysis** techniques such as recursion, amortized analysis;
 - A collection of **core** algorithms and data structures to solve a number problems from various computer science areas: algorithms for external memory, multiple-threaded algorithms, advanced graph algorithms, heuristic search and geometric calculations;
 - There will also enter into one or more **optional subjects** in advanced algorithms, including, but not limited to: *approximate algorithms*, randomized algorithms, search for text, linear programming and number theoretic algorithms such as crypto-systems.

Course Content (1)



- Lecture 1: Dynamic programming (CLRS 15)
 - Principles of DP.
 - Examples: Edit distance, activity selection.
- Lecture 2: All-pairs shortest paths (CLRS 25)
 - Distance matrix and predecessor matrix.
 - Floyd-Warshall algorithm, which uses DP.
- Lecture 3: Network flow algorithms (CLRS 26)
 - Formalize flow networks, flows, maximum-flow.
 - Ford-Fulkerson algorithms.
- Lecture 4: Greedy algorithm (CLRS 16)
 - Ideas/principles of greedy algorithm.
 - Examples: Activity selection, Huffman coding.

Course content (2)



- Lecture 5: Amortized analysis (CLRS 17)
 - Understand amortized analysis, difference from average-case analysis.
 - Aggregated analysis, accounting method, potential method.
- Lectures 6, 7, 8: Computational geometry (CLRS 33 + additional references)
 - Basic geometric operations in 2D, e.g., two line segments intersecting, orientation of two line segments?
 - Sweeping algorithms.
 - Graham's scan and Jarvis's march algorithm for convex hull.
 - Divide-and-conquer algorithm to find the closet pair of points in a set of points.
 - Range searching in d-dimensional space: kd-tree and range tree.

Course Content (3)



- Lecture 9: External-memory algorithms and data structures (CLRS 18 + additional references)
 - Balanced search trees, e.g., *B-trees* and *R-trees*.
 - External memory sorting: multi-way merge-sort algorithm.
- Lecture 10: Multi-threaded algorithms (CLRS 27)
 - Concurrency keywords: parallel, spawn, sync.
 - MT Fibonacci number computation, MT merge sort.
- Lecture 11 & 12: Algorithms for NP-complete problems (CLRS 35)
 - Approximation algorithms
 - Backtracking and branch-and-bound
 - Examples using the algorithms.
 - ◆ Vertex cover, traveling sales man.

Tips



- Your feedback is always welcome.
 - Esp. when you get confused.
 - Send me an email or drop by my office.
- Participate in every lecture.
- Play actively in every exercise class and self-study exercise session.
 - Exercises prepare you for the final exam!
 - Make sure you understand all exercises by YOURSELF after your group work in each exercise class.
- Check the course page frequently.

Agenda



- Introduction
- Dynamic Programming (DP)
 - To understand the principles of dynamic programming.
 - To understand the DP algorithm for edit distance.
 - To be able to apply the DP algorithm design technique.

Recall algorithm design techniques from AD1

- Algorithm design techniques so far:
 - Brute-force algorithms
 - ◆ Linear search
 - Incremental algorithms
 - ◆ Insertion sort
 - Algorithms that use ADTs (implemented using efficient data structures)
 - ◆ Heap sort
 - Divide-and-conquer algorithms
 - ◆ Merge sort, quick sort.
 - Dynamic-programming
 - ◆ Rod cutting
 - ◆ Top-down with memoization.
 - ◆ Bottom-up method.

Divide and Conquer



- If the problem size is small enough to solve it in a straightforward manner, solve it.
- Otherwise, meaning that the input size is too large to deal with in a straightforward manner, do the following
 - **Divide:** Divide the problem into two or more *disjoint* sub-problems.
 - **Conquer:** Use divide-and-conquer recursively to solve the sub-problems.
 - **Combine:** Take the solutions to the sub-problems and combine these solutions into a solution for the original problem.

Merge Sort



```
Merge-Sort(A, p, r)
  if p < r then
    q ← (p+r)/2
    Merge-Sort(A, p, q)
    Merge-Sort(A, q+1, r)
    Merge(A, p, q, r)
```

- Mini-quiz: do you still recall the **recurrence** of merge sort?
 - A **recurrence** is an equation or inequality that describes a function in terms of its value on smaller inputs.
- $T(n) =$
 - $\Theta(1)$ if $n=1$
 - $2T(n/2) + \Theta(n)$ if $n>1$

Dynamic programming



- A powerful technique to solve ***optimization problems***.
- An optimization problem can have many possible solutions, each solution has a value, and we wish to find a solution with the optimal (i.e., minimum or maximum) value.
- An algorithm should compute *the* **optimal value** plus, if needed, *an* **optimal solution**.


Rod cutting




- A steel rod of length n should be cut and sold in pieces.
- Pieces sold only in integer sizes according to a price table $P[1..n]$.
- *Goal*: cut up the rod to maximize profit.

Length	1	2	3	4	5	6	7
Price	4	5	13	16	23	24	27

Max profit (optimal value): 31
Optimal cut (optimal solution): 1, 1, 5

Price:  $23 + 5 = 28$


Price:  $16 + 13 = 29$

Price:  $13 + 13 + 4 = 30$

Two key characteristics of DP



- Overlapping sub-problems
 - Sub-problems share sub-sub-problems.
 - A divide-and-conquer algorithm does more work than necessary, as it needs to repeatedly solve the common sub-sub-problems.
- Optimal substructure
 - The optimal solution to a problem incorporates optimal solutions to sub-problems.
 - Un-weighted shortest path (YES)
 - ◆ Shortest path $\mathbf{A} = \langle q, r, t \rangle$ from q to t .
 - ◆ Sub-paths of \mathbf{A} , $\langle q, r \rangle$ and $\langle r, t \rangle$, are also the shortest paths.
 - Un-weighted longest simple path. (NO)
 - ◆ Longest path $\mathbf{B} = \langle q, r, t \rangle$ from q to t .
 - ◆ Sub-paths of \mathbf{B} , $\langle q, r \rangle$ and $\langle r, t \rangle$, may not be the longest paths.



$\langle q, s, t, r \rangle$
 $\langle r, q, s, t \rangle$

Two approaches of DP



- Top-down with memoization
 - Solve each sub-problem only **once** and **store** the answers to the solved sub-problems in a table.
 - Next time, when you need to solve a solved sub-problem, just look up the table to get the answer.
- Bottom-up without recursion.
 - Depending on some natural notion on the **size** of a sub-problem.
 - Solving any particular sub-problem depends only on solving **smaller** sub-problems.
 - Sort the sub-problems by size and solve them in size order, smallest first. And save the solutions.
- Pros and cons
 - Both should have the same asymptotic running time.
 - If all sub-problems must be solved, memoization (recursion) is usually slower (by a constant factor) than Bottom-up (loops).
 - If not all sub-problems need to be solved, memoization only solves the necessary ones.

Structure of DP



- Construction:
 - *What are the sub-problems?*
 - *Which choices have to be considered in order to solve a sub-problem?*
 - *How are the trivial sub-problems solved?*
 - *Write a memoized version of the algorithm or in which order do we have to solve the sub-problems (bottom-up)*
 - *Remember the (optimal) choices made*
 - *Use the remembered choices to construct a solution*
- Analysis:
 - *How many different sub-problems are there in total?*
 - *How many choices have to be considered when solving each sub-problem?*

Constructing
a solution

Recurrence for
the optimal value

Agenda



- Introduction
- Dynamic Programming
 - To understand the principles of dynamic programming.
 - To understand the DP algorithm for edit distance.
 - To be able to apply the DP algorithm design technique.

Edit distance



- Problem definition:
 - Two strings: $s[1..m]$ and $t[1..n]$
 - Find **edit distance** $dist(s, t)$ between the two input strings s and t .
 - ◆ The smallest number of **edit operations** that turns s into t .
 - Edit operations:
 - ◆ **Replace** one letter with another letter
 - ◆ **Delete** one letter
 - ◆ **Insert** one letter
- Example: let's turn “ghost” to “house”
 - **ghost** delete **g**
 - **host** insert **u**
 - **houst** replace **t** by **e**
 - **house**

Two cases



- The last letters in s and t are different, e.g., s=milk t=windy
 - Option 1: Replace k by y, $\text{dist}(s, t) = \text{dist}(\text{mil}, \text{wind}) + 1$
 - Option 2: Delete k, $\text{dist}(s, t) = \text{dist}(\text{mil}, \text{windy}) + 1$
 - Option 3: Insert y in the end of s, $\text{dist}(s, t) = \text{dist}(\text{milk}, \text{wind}) + 1$
 - $\text{dist}(s, t) = \min (\text{dist}(\text{mil}, \text{wind}) + 1, \text{dist}(\text{mil}, \text{windy}) + 1, \text{dist}(\text{milk}, \text{wind}) + 1)$
- The last letters in s and t are the same, e.g., s=milk t=link
 - Option 1: Keep k, $\text{dist}(s, t) = \text{dist}(\text{mil}, \text{lin})$
 - Option 2: Delete k, $\text{dist}(s, t) = \text{dist}(\text{mil}, \text{link}) + 1$
 - Option 3: Insert k in the end, $\text{dist}(s, t) = \text{dist}(\text{milk}, \text{lin}) + 1$
 - $\text{dist}(s, t) = \min (\text{dist}(\text{mil}, \text{lin}), \text{dist}(\text{mil}, \text{link}) + 1, \text{dist}(\text{milk}, \text{lin}) + 1)$
- Optimal sub-structure for edit distance?
 - YES! The optimal solution to a problem incorporates optimal solutions to sub-problems.
 - Formal proof: see Lecture Material on Moodle.

Sub-problems



- Sub-problem:
 - $d_{i,j} = \text{dist}(s[1..i], t[1..j])$
- Then $\text{dist}(s, t) = d_{m,n}$
- Let's look at the last symbol: $s[i]$ and $t[j]$. There are three options, do whatever is the cheapest:
- **Option 1:**
 - If $s[i] = t[j]$, then turn $s[1..i-1]$ to $t[1..j-1]$
 - ◆ milk, link: $d_{i,j} = d_{i-1,j-1}$
 - **Else replace** $s[i]$ by $t[j]$ and turn $s[1..i-1]$ to $t[1..j-1]$
 - ◆ milk, windy; **mily**, **windy**: $d_{i,j} = 1 + d_{i-1,j-1}$
- **Option 2: Delete** $s[i]$ and turn $s[1..i-1]$ to $t[1..j]$
 - milk, windy; **mil**, **windy**: $d_{i,j} = 1 + d_{i-1,j}$
- **Option 3: Insert** $t[j]$ at the end of $s[1..i]$ and turn $s[1..i]$ to $t[1..j-1]$
 - milk, windy; **milky**, **windy**: $d_{i,j} = 1 + d_{i,j-1}$

Recurrence, optimal substructure



$$d_{i,j} = \min \left\{ \begin{array}{ll} d_{i-1,j-1} + \begin{cases} 0 & \text{if } s[i] = t[j] \\ 1 & \text{else} \end{cases} & \begin{array}{l} \text{Do nothing} \\ \text{replace } s[i] \text{ by } t[j] \\ \text{delete } s[i] \end{array} \\ d_{i-1,j} + 1 & \\ d_{i,j-1} + 1 & \text{insert } t[j] \text{ at the end of } s[i] \end{array} \right.$$

- *How do we solve trivial sub-problems?*
 - To turn empty string to $t[1..j]$, do j **inserts**
 - To turn $s[1..i]$ to empty string, do i **deletes**
- Do the sub-problems overlap each other?

DP Algorithm, memoization



EditDistance(s[1..m], t[1..n])

```
01 for i = 0 to m do
02     for j = 0 to n do
03         dist[i, j] = ∞
04 return EditDistR(s, t, m, n)
```

Initialization

EditDistR(s, t, i, j)

```
01 if dist[i, j] == ∞ then
```

```
02     if j == 0 then dist[i, j] = i
```

```
03     else if i == 0 then dist[i, j] = j
```

```
04     else
```

```
05         if s[i] == t[j] then
```

```
06             dist[i, j] = min(EditDistR(s, t, i-1, j-1),
```

delete s [i]

insert t [j]

Replace

```
07             else
```

```
08                 dist[i, j] = 1 + min(EditDistR(s, t, i-1, j-1),
```

delete s [i]

insert t [j]

```
09 return dist[i, j]
```

Time Complexity



- Analysis
 - *How many different sub-problems are there in total?*
 - ◆ $n*m$
 - *How many choices have to be considered when solving each sub-problem?*
 - ◆ 3 (copy/replace, insert, and delete)
 - *Thus, $\Theta(nm)$*
 - If we solve editor distance in a naïve D&C manner, what is the complexity?
 - ◆ *Exponential runtime.*
 - ◆ See Moodle for full computations.

DP Algorithm, bottom-up



EditDistance(s[1..m], t[1..n])

```
01 for i = 0 to m do dist[i,0] = i
02 for j = 0 to n do dist[0,j] = j
03 for i = 1 to m do
04     for j = 1 to n do
05         if s[i] = t[j] then
06             dist[i,j] = min(dist[i-1,j-1], dist[i-1,j]+1,
                                dist[i,j-1]+1)
07         else
08             dist[i,j] = 1 + min(dist[i-1,j-1], dist[i-1,j],
                                dist[i,j-1])
09 return dist[m,n]
```

*Trivial sub-problems
i deletes and j inserts*

*Fills in entries in the
(m+1)*(n+1) matrix in
row-major order.*

- *What is the running time of this algorithm?*
- *How do we modify it to remember the edit operations?*

s="GO" t="LOG"



- m=2, n=3, we have a 3*4 matrix to fill in.

		j						
		0	1	L	2	O	3	G
i	0	0	1		2		3	
	1	G	1	1	2		2	
	2	O	2	2	1		2	

EditDistance(s[1..m], t[1..n])

01 **for** i = 0 **to** m **do** dist[i,0] = i

02 **for** j = 0 **to** n **do** dist[0,j] = j

03 **for** i = 1 **to** m **do**

04 **for** j = 1 **to** n **do**

05 **if** s[i] = t[j] **then**

06 dist[i,j] = min(dist[i-1,j-1], dist[i-1,j]+1,
 dist[i,j-1]+1)

07 **else**

08 dist[i,j] = 1 + min(dist[i-1,j-1], dist[i-1,j],
 dist[i,j-1])

09 **return** dist[m,n]

*Trivial sub-problems
i deletes and j inserts*

*Fills in entries in the
(m+1)*(n+1) matrix in
row-major order.*

Remember edit operations



- $m=2, n=3$, we have a 3×4 matrix to fill in.

0	1	L	2	O	3	G
---	---	---	---	---	---	---

0		0	1 I	2 I	3 I
1	G	1 D	1 R	2 R, I	2 C
2	O	2 D	2 R, D	1 C	2 I

```

EditDistance(s[1..m], t[1..n])
01 for i = 0 to m do dist[i,0] = i
02 for j = 0 to n do dist[0,j] = j
03 for i = 1 to m do
04     for j = 1 to n do
05         if s[i] = t[j] then
06             dist[i,j] = min(dist[i-1,j-1], dist[i-1,j]+1,
                                dist[i,j-1]+1)
07         else
08             dist[i,j] = 1 + min(dist[i-1,j-1], dist[i-1,j],
                                dist[i,j-1])
09 return dist[m,n]
    
```

*Trivial sub-problems
i deletes and j inserts*

*Fills in entries in the
(m+1)*(n+1) matrix in
row-major order.*

Remember edit operations



- $m=2$, $n=3$, we have a 3×4 matrix to fill in.

0	1	2	3
	L	O	G

0	0	1	2	3
1	D	R	R, I	C
2	D	R, D	C	I

$i=1$ Replace "L"

$i=2$ Copy "O"

$i=2$ Insert "G"

GO

LO

LO

LOG

What is the running time of this algorithm?

$\Theta(nm)$

Mini quiz



- $s = \text{"GO"}$ $t = \text{"LOGG"}$
- Fill in the 3×5 matrix.
- Identify the edit distance and the corresponding edit operations.
- There may be more than one possible sequences of operations.

ILO of Lecture 1



- Dynamic Programming
 - To understand the principles of dynamic programming.
 - ◆ Overlapping sub-problems and optimal sub-structure.
 - ◆ Top-down with memoization and bottom-up.
 - To understand the DP algorithm for edit distance.
 - To be able to apply the DP algorithm design technique.

Lecture 2



- All-pairs shortest paths (dynamic programming)
 - To understand the adjacency matrix and the predecessor matrix, which are the representations of the input and output of most of the all-pairs shortest-path algorithms.
 - To understand how the dynamic programming principles play out in the Floyd-Warshall algorithm.