

Advanced Algorithms

Lecture 10

Multithreaded Algorithms

Tung Kieu

tungkvt@cs.aau.dk

Center for Data-intensive Systems

ILO of Lecture 10



- Multithreaded algorithms
 - to understand the model of dynamic multithreading, including **nested parallelism** and **parallel loops**;
 - to understand **work**, **span**, and **parallelism** — the concepts necessary for the analysis of multithreaded algorithms;
 - to understand and be able to analyze the multithreaded merge sort algorithm.

Agenda



- Background
- Nested parallelism
- Work, span, and parallelism
- Parallel loops
- Multithreaded merge sort

Parallel computers



- Computers with multiple processing units
 - Chip multiprocessors / Low price
 - ◆ A single multi-core chip has multiple
 - ◆ My laptop: Intel® Core™ i7-7700U P
 - Clusters / Intermediate price.
 - ◆ Individual computers that are connected
 - Supercomputers / High price.
 - ◆ Combination of custom architectures for the highest performance.
- Memory models
 - Shared memory: all processors can access any location in memory.
 - Distributed memory: where each processor has a private memory.
 - No agreement on a single architectural model for parallel computers so far.



Static vs. dynamic threading



- Static threading
 - Each thread maintains an associated program counter and executes code independently of the other threads.
 - Threads persist for the duration of a computation.
 - Directly using static threading is difficult and error-prone.
- Dynamic threading
 - Specify parallelism in applications without worrying implementation details.
 - A **concurrency platform's scheduler** does communication, load balancing, etc.
 - ***Nested parallelism and parallel loops.***

Agenda



- Background
- Nested parallelism
- Work, span, and parallelism
- Parallel loops
- Multithreaded merge sort

Fibonacci Numbers



- *Leonardo Fibonacci (1202):*
 - We have a rabbit in the beginning.
 - A rabbit starts producing offspring on the second generation after its birth and produces one child each generation.
 - How many rabbits will there be after n generations?

$$F(1)=1$$



Fibonacci Numbers (2)



- $F(n) = F(n-1) + F(n-2)$
- $F(0) = 0, F(1) = 1$
- 0, 1, 1, 2, 3, 5, 8, 13, 21, 34 ...

FIB(n)

1 **if** $n \leq 1$

2 **return** n


3 **else** $x = \text{FIB}(n - 1)$

4 $y = \text{FIB}(n - 2)$

5 **return** $x + y$


- Straightforward recursive procedure is **very very slow!**
- Why? How slow? **$O(2^n)$**
- Let's draw the recursion tree.

Fibonacci Numbers (3)



- We keep calculating the same value over and over!
 - Sub-problems are overlapping – they share sub-sub-problems.
- Do you still remember how we can avoid calculating the same sub-problems?

Fibonacci Numbers (3)



- We keep calculating the same sub-problems
 - Sub-problems are overlapping
- Do you still remember how we can avoid calculating the same sub-problems?

Dynamic Programming
Memoization (Top Down)
Tabulation (Bottom Up)

Fibonacci Numbers (4)



- Recurrence

- $T(n) = T(n-1) + T(n-2) + \Theta(1)$ Master theorem

- Since $T(n-1) \geq T(n-2)$, we have $T(n) \geq 2T(n-2) + a$

- Solving the recurrence using the repeated substitution method.

- $T(n) = 2T(n-2) + a \rightarrow T(n-2) = 2T(n-4) + a$

- $T(n) = 2^2T(n-4) + (2+1)a \rightarrow T(n-4) = 2T(n-6) + a$

- $T(n) = 2^3T(n-6) + (2^2+2+1)a$

- $T(n) = 2^iT(n-2*i) + (2^{i-1} + \dots + 2 + 1)a = 2^iT(n-2*i) + a \sum_{k=0}^{i-1} 2^k$

- When $i=n/2$, $T(n) = 2^{n/2}T(0) + a * 2^{n/2} = (a+1) 2^{n/2}$

$$\sum_{k=0}^n x^k = \frac{x^{n+1} - 1}{x - 1}$$

- We get


- $T(n) \geq (a+1) 2^{n/2} \approx (a+1) 1.4^n$

- Running time is **exponential!**

Multithreaded version



- What if we can do the two recursive calls in parallel?
- Using the so-called *nested parallelism*
 - Call a procedure. Don't wait for it to return and go to next step.
 - Call Fib(5), we do not need to wait for Fib(5) returns results and we can call Fib(4).



Concurrency keywords



- Spawn
 - If the keyword **spawn** precedes a procedure call, it indicates a *nested parallelism*.
 - No need to wait for the procedure with keyword **spawn**.
 - ◆ We do not wait P-FIB($n-1$) to complete and we can execute P-FIB($n-2$) on another processing unit.
- Sync
 - Must wait for all spawned procedures to complete before going to the statement after **sync**.
 - ◆ Before “return $x+y$ ”, we must wait P-FIB($n-1$) to complete.

P-FIB(n)

```
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{spawn P-FIB}(n - 1)$ 
4       $y = \text{P-FIB}(n - 2)$ 
5      sync
6      return  $x + y$ 
```

FIB(n)

```
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{FIB}(n - 1)$ 
4       $y = \text{FIB}(n - 2)$ 
5      return  $x + y$ 
```

Analysis



- Recurrence:
 - $T(n) = \max(T(n-1), T(n-2)) + c = T(n-1) + c$
 - Can you try to solve this recurrence?
- Solving the recurrence, we have $T(n) = \Theta(n)$.
- The serial version is **exponential** whereas the multithreaded version is only **linear**.

$$T(n) = T(n-1) + c$$

$$T(n-1) = T(n-2) + c$$

$$\Rightarrow T(n) = (T(n-2) + c) + c = T(n-2) + 2c$$

$$T(n-2) = T(n-3) + c$$

$$\Rightarrow T(n) = T(n-3) + 3c$$

$$\Rightarrow \dots$$

$$\Rightarrow T(n) = T(0) + nc$$

Agenda




- Background
- Nested parallelism
- Work, span, and parallelism
- Parallel loops
- Multithreaded merge sort

Work, Span, Parallelism



- Three main concepts (informally):
 - *Work* – the running time on a machine with one-processor (T_1).
 - ◆ Fibonacci: $\Theta(\varphi^n)$
 - *Span* – the running time on a machine with infinite processors (T_∞).
 - ◆ Fibonacci: $\Theta(n)$
 - *Parallelism* = *Work*/*Span* – how many processors on average are used by the algorithm.
 - ◆ Fibonacci: $\Theta(\varphi^n / n)$
- More formally, work and span are defined using a computation DAG (directed acyclic graph).
 - Vertices are instructions or sets of instructions.
 - Edges represent dependencies between instructions.
 - ◆ An edge (u, v) means that u must execute before v .



Computation DAG



- Using an example of computing Fibonacci number of 4.


P-FIB(n)

```
1  if  $n \leq 1$ 
2    return  $n$ 
3  else  $x = \text{spawn P-FIB}(n - 1)$ 
4       $y = \text{P-FIB}(n - 2)$ 
5    sync
6    return  $x + y$ 
```

● Lines 1-3


● Lines 4-5

○ Line 6




Computation DAG


Edge(u, v) means that u must execute before v .



Spawned procedure



Called procedure



$P-FIB(n)$

```

1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{spawn } P-FIB(n - 1)$ 
4       $y = P-FIB(n - 2)$ 
5      sync
6      return  $x + y$ 
    
```

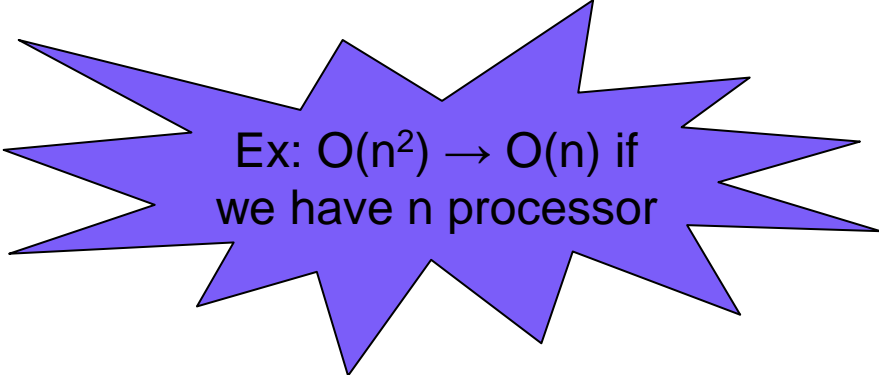
Work: number of vertices, 17

Span: the length of the longest path (critical path), 8

Work law and span law



- Notation
 - Work T_1
 - Span T_∞
 - Multithreaded computation on P processors: T_P
- Work law: $T_P \geq T_1 / P$
 - An ideal parallel computer with P processors can do at most P units of work.
- Span law: $T_P \geq T_\infty$
 - An ideal parallel computer with P processors cannot run any faster than a machine with unlimited number of processors.



Ex: $O(n^2) \rightarrow O(n)$ if
we have n processor

Speedup



- Speedup: T_1/T_P :

Work law

- How many times faster the computation on P processors than on a single processor.
- Recall the work law says that $T_P \geq T_1 / P$, so that the speedup must be smaller than or equal to P .
 - ◆ The speedup on a P -processor machine can **at most** be P .
 - ◆ If the speedup is P , we have **perfect linear speedup**.

Mini quiz



- Considering the case for computing P-Fib(4).
- We already know that the work $T_1 = 17$ and the span $T_\infty = 8$
- Consider the following setups, each setup corresponds to a machine with P processing units. Which one is the most likely setup to achieve the ***perfect linear speedup***?
- $P=2, P=3, P=4, P=5, P=10$

Mini quiz



- Considering the case for computing P-Fib(4). We already know that the work $T_1 = 17$ and the span $T_\infty = 8$
- Consider the following setups, each setup corresponds to a machine with P processing units. Which one is the most likely setup to achieve the **perfect linear speedup**?
- $P=2$, $P=3$, $P=4$, $P=5$, $P=10$
- Span law: $T_P \geq T_\infty$
 - $P=2$, having perfect linear speedup indicates $T_2 = 17/2 = 8.5 \geq T_\infty$, which is possible.
 - $P=3$, having perfect linear speedup indicates $T_3 = 17/3 = 5.7 \leq T_\infty$, which is impossible.
 - So for $P=4$, $P=5$, and for any P if P is greater than 2.

Parallelism and Slackness



- Parallelism: T_1 / T_∞ :
 - The **maximum possible speedup** that can be achieved on any number of processors.
 - Once the number of processors exceeds the parallelism, the computation cannot possibly achieve the perfect linear speedup.
- Slackness: $\text{Parallelism}/P = T_1 / PT_\infty$
 - Assume that P increases from 1 each time by 1.
 - Slackness drops from Parallelism to 1 and then from 1 to 0.
 - Slackness ≥ 1 : closer to perfect linear speedup.
 - Slackness < 1 : diverges further away from perfect linear speedup.
- Example: parallelism = 10.
 - $P=2$, slackness=5, perfect linear speedup is possible.
 - $P=10$, slackness=1, perfect linear speedup is possible but more difficult than when $P=2$
 - $P=20$, slackness=0.5, impossible to achieve perfect linear speedup, i.e., 20 speed up, due to the parallelism=10

Parallelism and Slackness



- Parallelism: T_1 / T_∞ :
 - The **maximum possible speedup** that can be achieved on any number of processors.
 - Once the number of processors exceeds the parallelism, the computation cannot possibly achieve the perfect linear speedup.
- Slackness: $\text{Parallelism}/P = T_1 / (P \cdot T_P)$
 - Assume that P increases.
 - Slackness drops from Parallelism.
 - Slackness ≥ 1 : closer to perfect linear speedup.
 - Slackness < 1 : diverges further away from perfect linear speedup.
- Example: parallelism = 10.
 - $P=2$, slackness=5, perfect linear speedup is possible.
 - $P=10$, slackness=1, perfect linear speedup is possible but more difficult than when $P=2$
 - $P=20$, slackness=0.5, impossible to achieve perfect linear speedup, i.e., 20 speed up, due to the parallelism=10

$$17/8$$

$$17/P8$$

$$P=1 \rightarrow 17/8$$

$$P=2 \rightarrow 17/16$$

$$P=3 \rightarrow 17/24$$

To Summarize



Notation	Meaning
T_1	Work, the running time on a machine with one processor.
T_∞	Span, the running time on a machine with infinite processors.
T_P	The running time on a machine with P processors.
$T_P \geq T_1 / P$	Work law
$T_P \geq T_\infty$	Span law
T_1 / T_P	Speedup. Speedup must be $\leq P$ according to the work law. When speedup is equal to P, it achieves perfect speed up .
T_1 / T_∞	Parallelism. The maximum possible speedup that can be achieved on any number of processors
T_1 / PT_∞	Slackness = Parallelism/P. The larger the slackness, the more likely to achieve perfect speed up. When slackness is less than 1, it is impossible to achieve perfect speed up.

Analyzing multithreaded algorithms




- We have two sub-computations A and B.
- If they are joined in series:



$$\text{Work: } T_1(A \cup B) = T_1(A) + T_1(B)$$

$$\text{Span: } T_\infty(A \cup B) = T_\infty(A) + T_\infty(B)$$

- If they are joined in parallel:



$$\text{Work: } T_1(A \cup B) = T_1(A) + T_1(B)$$

$$\text{Span: } T_\infty(A \cup B) = \max(T_\infty(A), T_\infty(B))$$

Analyzing multithreaded Fibonacci



- $T_{\infty}(n) = \max(T_{\infty}(n-1), T_{\infty}(n-2)) + c$
- $T_{\infty}(n) = T_{\infty}(n-1) + c$ $T_{\infty}(n-2) \leq T_{\infty}(n-1)$
- ... $T_{\infty}(n-1) = \max(T_{\infty}(n-2) + T_{\infty}(n-3)) + c$
- $T_{\infty}(n) = T_{\infty}(0) + nc$
- $T_{\infty}(n) = \Theta(n)$.
- Parallelism of P-FIB(n) is $T_1(n) / T_{\infty}(n) = \Theta(\varphi^n / n)$
 - As n increases, the parallelism grows.

Agenda



- Background
- Nested parallelism
- Work, span, and parallelism
- Parallel loops
- Multithreaded merge sort

Concurrency keywords: parallel



- **Parallel** for parallel loops

- Just like a for loop but loop executions run concurrently

```
ArrayCopy (A, B)
```

```
01 parallel for i = 1 to sizeof(A) do
```

```
02     B[i] = A[i]
```

- Implemented in a divide-and-conquer manner by using spawn.

```
ArrayCopyRecursive (A, B, l, r)
```

```
01 if l > r then return
```

```
02 if l = r then B[l] = A[r]
```

```
03 else
```

```
04     q =  $\lceil (l+r)/2 \rceil$ 
```

```
05     spawn ArrayCopyRecursive (A, B, l, q-1)
```

```
06         ArrayCopyRecursive (A, B, q, r)
```

Here, we only make 1 spawn.

What if we make 2 spawns?

And what if we make n-1 spawns?

Span: $T_{\infty}(n) = \max(T_{\infty}(n/2), T_{\infty}(n/2)) = T_{\infty}(n/2) + c = \Theta(\lg n)$

Work: $T_1(n) = 2T_1(n/2) + c = \Theta(n)$

Why not spawn n-1 times?



- Creating a thread is not **free**, but takes some time.
 - Assume that it takes constant time **c**.
- Implemented in a divide-and-conquer manner by using one (or a constant number of) spawn(s).

```
ArrayCopyRecursive(A, B, l, r)
01 if l > r then return
02 if l = r then B[l] = A[r]
03 else
04   q =  $\lceil (l+r)/2 \rceil$ 
05   spawn ArrayCopyRecursive(A, B, l, q-1)
06         ArrayCopyRecursive(A, B, q, r)
```

- Using one spawn
 - Span: $T_{\infty}(n) = \max(T_{\infty}(n/2), T_{\infty}(n/2)) + c$
 - $T_{\infty}(n) = T_{\infty}(n/2) + c$
 - $\Theta(\lg n)$
 - Work: $T_1(n) = T_1(n/2) + T_1(n/2)$
 - $T_1(n) = 2T_1(n/2) + c$
 - $\Theta(n)$

Why not spawn $n-1$ times?



- Creating a thread is not **free**, but takes some time.
 - Assume that it takes constant time **c**.
- Using two spawns
 - Span: $T_{\infty}(n) = \max(T_{\infty}(n/3), T_{\infty}(n/3), T_{\infty}(n/3)) + 2c$
 - $T_{\infty}(n) = T_{\infty}(n/3) + 2c$
 - $\Theta(\lg n)$
 - Work: $T_1(n) = 3T_1(n/3) + 2c$
 - $\Theta(n)$

Why not doing n copies in parallel?



- Creating a thread is not **free** but takes some time.
 - Assume that it takes constant time c .
- If we make **$n-1$** spawn threads and each thread copies one element, then we have the following:
 - Span: $T_{\infty}(n) = \max(T_{\infty}(n/\mathbf{n}), \dots, T_{\infty}(n/\mathbf{n})) + (\mathbf{n-1}) * \mathbf{c} =$
 - $T_{\infty}(n) = a + (\mathbf{n-1}) * \mathbf{c}$
 - $\Theta(n)$
 - Work: $T_1(n) = nT_1(1) + (\mathbf{n-1}) * \mathbf{c}$
 - $\Theta(n)$

Why not doing n copies in parallel?



- Creating a thread is not **free** but takes some time.
 - Assume that it takes constant time c .
- Assume that we have P processing units, and we make $P-1$ spawn threads.
 - Span: $T_P(n) = \max(T_P(n/P), \dots, T_P(n/P)) + (P-1)*c$
 - $T_P(n) = T_P(n/P) + (P-1)*c$
 - $\Theta(P \lg_P n)$
 - Assume that $P=2$, $T_P(n) = \Theta(2 \lg_2 n) = \Theta(\lg n)$
 - Assume that $P=3$, $T_P(n) = \Theta(3 \lg_3 n) = \Theta(\lg n)$
 - Assume that $P=n$, $T_P(n) = \Theta(n \lg_n n) = \Theta(n)$

Race conditions



- Race conditions happen when
 - Two logically parallel instructions access the same memory location and at least one of the instructions performs a write.

RACE-EXAMPLE()

```
1  x = 0
2  parallel for i = 1 to 2
3      x = x + 1
4  print x
```

- How to avoid?
 - Need to make sure that the two threads are **independent** - no writing to memory that other threads are reading from.

Matrix-vector multiplication



A $n \times n$ matrix multiplies a $n \times 1$ vector. We get a $n \times 1$ vector.

$$y_i = \sum_{j=1}^n a_{ij} x_j \quad \begin{array}{l} i: \text{row} \\ j: \text{column} \end{array}$$

a_{11}	a_{12}	...	a_{18}
a_{21}	a_{22}	...	a_{28}
...
a_{81}	a_{82}	...	a_{88}

x_1
x_2
...
x_8

MAT-VEC(A, x)

```
1   $n = A.rows$ 
2  let  $y$  be a new vector of length  $n$ 
3  parallel for  $i = 1$  to  $n$ 
4       $y_i = 0$ 
5  parallel for  $i = 1$  to  $n$ 
6      for  $j = 1$  to  $n$ 
7           $y_i = y_i + a_{ij} x_j$ 
8  return  $y$ 
```


Can we make this also a parallel loop?

- The first parallel loop: same with ArrayCopy that we just saw.
- The second parallel loop: a bit complicated.
 - Still use D&C and Spawn.
 - Base case is the for loop with j .

Matrix-vector multiplication



A $n \times n$ matrix multiplies a $n \times 1$ vector. We get a $n \times 1$ vector.



i : row
 j : column

a_{11}	a_{12}	...	a_{18}
a_{21}	a_{22}	...	a_{28}
...
a_{81}	a_{82}	...	a_{88}

x_1
x_2
...
x_8

MAT-VEC(A, x)

```
1   $n = A.rows$ 
2  let  $y$  be a new vector of length  $n$ 
3  parallel for  $i = 1$  to  $n$ 
4       $y_i = 0$ 
5  parallel for  $i = 1$  to  $n$ 
6      for  $j = 1$  to  $n$ 
7           $y_i = y_i + a_{ij}x_j$ 
8  return  $y$ 
```

Can we make this also a parallel loop?

- The first parallel loop: same with ArrayCopy that we just saw.
- The second parallel loop: a bit complicated.
 - Still use D&C and Spawn.
 - Base case is the for loop with j .

The second parallel loop

```

5  parallel for  $i = 1$  to  $n$ 
6      for  $j = 1$  to  $n$ 
7           $y_i = y_i + a_{ij}x_j$ 

```

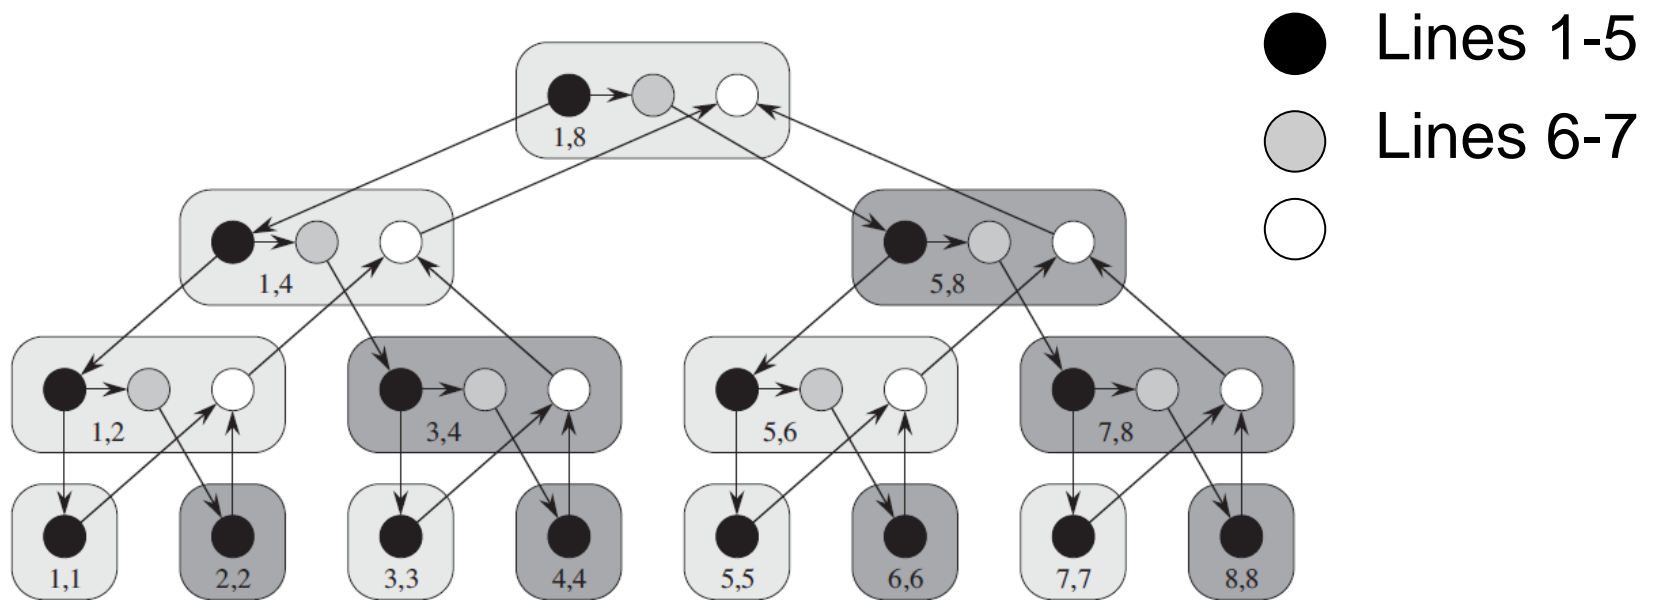
MAT-VEC-MAIN-LOOP(A, x, y, n, i, i')

```

1  if  $i == i'$ 
2      for  $j = 1$  to  $n$ 
3           $y_i = y_i + a_{ij}x_j$ 
4  else  $mid = \lfloor (i + i')/2 \rfloor$ 
5      spawn MAT-VEC-MAIN-LOOP( $A, x, y, n, i, mid$ )
6      MAT-VEC-MAIN-LOOP( $A, x, y, n, mid + 1, i'$ )
7      sync

```

A: input matrix.
x: input vector.
y: output vector.
n: the dimension of the vector.
i and i': start and end positions.



Analysis



- Recurrence for the span.


- $T_{\infty}(n) = \begin{cases} \Theta(n), & \text{if } n=1 \\ \max(T_{\infty}(n/2), T_{\infty}(n/2)) + c = T_{\infty}(n/2) + c, & \text{if } n>1 \end{cases}$
- We need to solve the recurrence.
 - ◆ $T_{\infty}(n) = T_{\infty}(n/2) + c$ $T_{\infty}(n/2) = T_{\infty}(n/4) + c$
 - ◆ $= T_{\infty}(n/2^2) + 2c$ $T_{\infty}(n/4) = T_{\infty}(n/8) + c$
 - ◆ $= T_{\infty}(n/2^3) + 3c$
 - ◆ $= T_{\infty}(n/2^i) + i \cdot c$
 - ◆ In order to get to the base case, let $2^i = n$, so $i = \lg n$.
 - ◆ $T_{\infty}(n) = T_{\infty}(n/2^i) + i \cdot c$
 - ◆ $= T_{\infty}(1) + \lg n \cdot c$
 - ◆ $= \Theta(n) + \Theta(\lg n)$
 - ◆ $= \Theta(n)$
- So the span is $\Theta(n)$.

Analysis




- Recurrence for the work.
 - $T_1(n) = \begin{cases} \Theta(n), & \text{if } n=1 \\ 2T_1(n/2) + c, & \text{if } n>1 \end{cases}$
 - The work is $\Theta(n^2)$ due to the nested loops in 5-7
- The parallelism $\Theta(n^2) / \Theta(n) = \Theta(n)$

Analysis of parallel loops

iteration_∞(i): 

- $T_{\infty}(n) = \Theta(\lg n) + \max_{1 \leq i \leq n} \text{iteration}_{\infty}(i)$

the span of
the i-th iteration in
the parallel loop.



Height of the
recursion tree
is $\Theta(\lg n)$

- Arraycopy $\max_{1 \leq i \leq n} \text{iteration}_{\infty}(i)$
 - $\max_{1 \leq i \leq n} \text{iteration}_{\infty}(i) = \Theta(1)$, constant time, because of no iteration.
 - $T_{\infty}(n) = \Theta(\lg n) + \Theta(1) = \Theta(\lg n)$
- Matrix-vector multiplication
 - $\max_{1 \leq i \leq n} \text{iteration}_{\infty}(i) = \Theta(n)$
 - $T_{\infty}(n) = \Theta(\lg n) + \Theta(n) = \Theta(n)$

Agenda



- Background
- Nested parallelism
- Work, span, and parallelism
- Parallel loops
- Multithreaded merge sort

Merge sort



```
Merge-Sort(A, p, r)
01 if p < r then
02     q = ⌊(p+r)/2⌋
03     Merge-Sort(A, p, q)
04     Merge-Sort(A, q+1, r)
05     Merge(A, p, q, r)
```

} *Divide*
 } *Conquer*
 } *Combine*

$$T(n) = aT(n/b) + f(n)$$

$$a=2, b=2, f(n) = n$$

$$\rightarrow n^{\log_b a} = n^{\log_2 2} = n == f(n)$$

$$T(n) = \Theta(n \lg n)$$

- Run-time?

- $T(n) = 2T(n/2) + \Theta(n)$
- $\Theta(n \lg n)$

Multithreaded merge sort



```
Merge-Sort'(A, p, r)
01 if p < r then
02   q = ⌊(p+r)/2⌋
03   spawn Merge-Sort'(A, p, q)
04   Merge-Sort'(A, q+1, r)
05   sync
06   Merge(A, p, q, r)
```

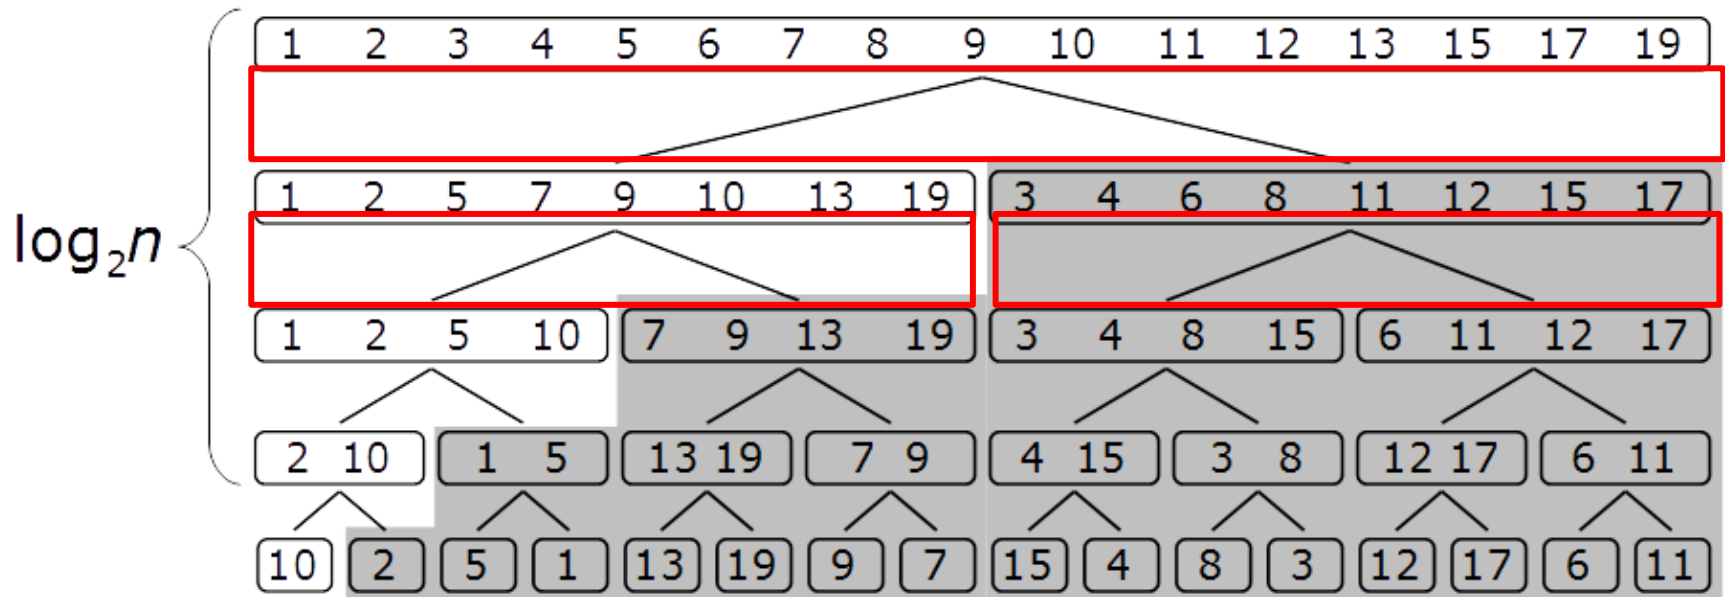
- Work: use W to denote T_1
 - $W(n) = 2W(n/2) + \Theta(n)$
 - $= \Theta(n \lg n)$, the same with serialized merge sort.
- Span: use S to denote T_∞
 - $S(n) = \max(S(n/2), S(n/2)) + \Theta(n)$
 - $= S(n/2) + \Theta(n)$
 - $= \Theta(n)$
- Parallelism
 - $W(n)/S(n) = \Theta(n \lg n) / \Theta(n) = \Theta(\lg n)$.

The parallelism is rather low.
Can we do better, i.e., higher parallelism?
If yes, which part can we further improve?

Merge



- The bottleneck: the Merge() function is very serial.



- At the top level, only one processor does $\Theta(n)$ work in serial!
- At the second level, only two processors do $\Theta(n)$ work.
- ...
- Can we make Merge more parallel?


Multithreaded merge



Main idea – make Merge() divide-and-conquer and use nested parallelism.

P-Merge

Input



1. Get middle element, pivot


2. Do binary search

Output


1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16

3. Place the pivot to the output, do 4 and 5 in *parallel*
4. Merge left halves, place the result to the left of pivot
5. Merge right halves, place the result to the right of pivot

P-Merge



P-Merge



P-Merge

T: input array.

$p_1 \dots r_1$ and $p_2 \dots r_2$ in T: two sorted sub-arrays

A: output array.

P-MERGE($T, p_1, r_1, p_2, r_2, A, p_3$) p_3 : starting position in A.

```
1   $n_1 = r_1 - p_1 + 1$ 
2   $n_2 = r_2 - p_2 + 1$ 
3  if  $n_1 < n_2$                                 // ensure that  $n_1 \geq n_2$ 
4      exchange  $p_1$  with  $p_2$ 
5      exchange  $r_1$  with  $r_2$ 
6      exchange  $n_1$  with  $n_2$ 
7  if  $n_1 == 0$                                 // both empty?
8      return
9  else  $q_1 = \lfloor (p_1 + r_1) / 2 \rfloor$ 
10      $q_2 = \text{BINARY-SEARCH}(T[q_1], T, p_2, r_2)$ 
11      $q_3 = p_3 + (q_1 - p_1) + (q_2 - p_2)$ 
12      $A[q_3] = T[q_1]$ 
13     spawn P-MERGE( $T, p_1, q_1 - 1, p_2, q_2 - 1, A, p_3$ )
14     P-MERGE( $T, q_1 + 1, r_1, q_2, r_2, A, q_3 + 1$ )
15     sync
```

Why do we need to make sure $n_1 \geq n_2$?

q_1 : the middle element's index.

$T[q_1]$: the pivot element.

Binary search for $T[q_1]$ from p_2 to r_2 elements in T.

Place the pivot in the right place in A


Merge the left and right halves in parallel.

P-Merge



- Why do we need to make sure $n_1 \geq n_2$?
- It makes sure that the maximum number of elements in either of the divide-and-conquer merge calls is $3n/4$

- $n_2 = (n_2 + n_2)/2 \leq (n_1 +$
- In the worst case:



If we select any of these elements in red as a pivot, then both L and R will have size $\leq \frac{3}{4}n$.

- We have $\lfloor n_1/2 \rfloor$ elements
 - We have n_2 elements
 - ◆ $\lfloor n_1/2 \rfloor + n_2$
 - ◆ $\leq n_1/2 + n_2 = n_1/2 + n_2/2 + n_2/2$
 - ◆ $= (n_1 + n_2)/2 + n_2/2$
 - ◆ $= n/2 + n_2/2$
 - ◆ $\leq n/2 + n/4$
 - ◆ $= 3n/4$

Running time for binary search.

- Span: $S(n) = \max(S(\alpha n), S((1 - \alpha)n)) + \Theta(\lg n)$
- $= S(3n/4) + \Theta(\lg n)$. See exercise 4.6-2, CLRS.
- $S(n) = \Theta(\lg^2 n)$ Extension to Master Method.

Multithreaded Merge Sort with P-Merge

P-MERGE-SORT(A, p, r, B, s)

```
1   $n = r - p + 1$ 
2  if  $n == 1$ 
3       $B[s] = A[p]$ 
4  else let  $T[1..n]$  be a new array
5       $q = \lfloor (p + r) / 2 \rfloor$ 
6       $q' = q - p + 1$ 
7      spawn P-MERGE-SORT( $A, p, q, T, 1$ )
8      P-MERGE-SORT( $A, q + 1, r, T, q' + 1$ )
9      sync
10     P-MERGE( $T, 1, q', q' + 1, n, B, s$ )
```

- Work: $W(n) = \Theta(n \lg n)$
- Span: $S(n) = \max(S(n/2) + S(n/2)) + \Theta(\lg^2 n) = S(n/2) + \Theta(\lg^2 n)$
 - $\Theta(\lg^3 n)$
- Parallelism: $\Theta(n \lg n) / \Theta(\lg^3 n) = \Theta(n / \lg^2 n)$
 - This is a better (higher) parallelism compared to $\Theta(\lg n)$.

P-Merge-Sort P-Merge

To summarize



- Work: $\Theta(n \lg n)$

	Merge Procedure	Span	Parallelism
Naïve merge	$\Theta(n)$	$\Theta(n)$	$\Theta(\lg n)$
P-Merge	$\Theta(\lg^2 n)$	$\Theta(\lg^3 n)$	$\Theta(n / \lg^2 n)$

Goal of the multi-threaded algorithm design



- Goal of the multi-threaded algorithm design – *increase parallelism*.
 - $Parallelism = work / span$
 - Usually achieved by decreasing span
 - ◆ MergeSort without P-Merge and with P-Merge
 - ◆ $\Theta(n)$ vs. $\Theta(\lg^3 n)$
 - It may pay off to slightly increase work, if span can be decreased significantly (in practice, relevant for highly parallel systems, such as supercomputers, GPUs)

ILO of Lecture 10



- Multithreaded algorithms
 - to understand the model of dynamic multithreading, including **nested parallelism** and **parallel loops**;
 - to understand **work, span,** and **parallelism** — the concepts necessary for the analysis of multithreaded algorithms;
 - to understand and be able to analyze the multithreaded merge sort algorithm.